

Dharmik Patel - dsp187  
Dhruv Chaudhry - dc1533  
ilab: cp.cs.rutgers.edu

#### Makefile:

- make rufs: run this command to compile the rufs.c
- make check\_mt: run this command to check if the DISKFILE is mounted
- make remove\_mt: run this command to remove the mount. Helpfull when rufs exits without calling rufs\_destroy()
- make run\_fuse: run this command to run our custom file System
- make clean: remove all compiled files AND the DISKFILE. (erases our 'HDD')
- our mount is at /tmp/dsp187/mountdir

#### Benchmarks:

- simple\_test.c:
  - Run time: 1.4 milliseconds
  - Number of data blocks used: 123 dblocks
- test\_cases.c:
  - Run time: 2.1 milliseconds
  - Number of data blocks used: 123 dblocks
- Notes:
  - Run time was calculated using clock() from time.h in milli seconds
  - Number of data blocks used, does not include superblocks, inode blocks, and bitmap blocks. It is only the data blocks.
  - just printed out the number of set bits in dblock\_bm when rufs\_destroy was called

#### Implementation:

- There were 4 sets of functions that we implemented. The entry functions, the helper functions, dir functions and file functions
- Entry Functions:
  - rufs\_mkfs:
    - this was called every time we were starting from an unformatted DISKFILE
    - it setup the root dir, with a single child dir '.'
    - we also setup the bitmaps with calloc to make sure they were all zero
  - rufs\_init:
    - the entry point of FUSE. if the disk was already formatted, we just loaded the superblock into mem.
    - this also served as a test bed for our helper functions before we implemented the dir and file functions
  - rufs\_destroy:
    - if the program exited with ^C, this function would get called, and free all the mem, and close the DISKFILE
    - we also printed the amount dblocks used at this point. see section above on Benchmarks.
  - rufs\_getattr:

- this would be called a lot of times by FUSE
- if a file/dir was not found, then it would return -ENOENT and set errno to ENOENT.
- Note on perm for REG files:
  - Abhilash, TA, said we can ignore the perms while creating a file for this project. so I just set them to rw-r--r--. the benchmarks create files with rw-rw-rw-. But that is not important for this project

- Helper Functions:

- my\_print:

- print only when DEBUG was set to 1, except the value of DEBUG did not matter in my\_print\_always().

- these set of functions were a wrapper around printf, and helped us debug the code effectively with the different color outputs

- bitmap functions:

- using the already provided bitmap functions, we implemented get\_avail\_ino() and get\_avail\_blkno().

- every time we would run these 2 functions, we would make sure to do bio\_read() and bio\_write(), to ensure our changes were persistent on disk

- readi() and writei():

- used to write and read inodes

- made sure to calculate the proper offset needed since there could be multiple inodes in 1 inode block

- bio\_read() and bio\_write():

- although these implementations were provided, we made sure to that every time we used, we would:

- 1. make sure the buffer was of size BLOCK\_SIZE

- 2. if we were filling the buffer in from scratch, and not from bio\_read(), we would calloc the mem used by buffer

- 3. when writing data blocks that were stored in the direct pointers, we would always add 'sb->d\_start\_blk'

- dir\_add():

- used to add a file or subdir to a given parent\_dir

- we would make sure to update the parent\_dir hardlink counter

- we would also calloc a new block if needed to add a new dirent

- thoroughly tested by the Benchmarks

- dir\_find():

- used to check if a file or subdir was in a given parent\_dir

- make sure that the parent\_dir is a dir and not a file

- get\_node\_by\_path():

- we had a recursive approach to this

- it would find a inode related to the absolute path given

Once these were implemented, the rest of the functions were easy

- Dir Functions:

- rufs\_opendir:

- this would check if a given absolute path was a dir and if it existed

rufs\_readdir:

- used to list out all the dirents in a given dir
- used the filler function to put the dirnames

rufs\_mkdir:

- used to add a subdir at a given absolute path
- made sure that the dirname() was a dir and had enough space to add a dirent
- then we made a new dir with the normal subdirs '.' and '..' and made sure to link them

properly.

dot would link to basename dir

dotdot would link to parent\_dir

- File Functions:

rufs\_create:

- used to create a REG file
- did not allocate a data block yet, since the file size is 0 bytes

rufs\_open:

- this would check if a given absolute path was a reg file and if it existed

rufs\_write:

- write data to the data blocks
- made sure to use bio\_read() on the starting block that contained the offset, to make sure

all data before offset was preserved

- same for the block containing the end, so we don't write past offset+size
- all full blocks in the middle could just be overridden, so no need to waste bio\_read
- if new data blocks allocated, made sure to also update the inode

rufs\_read:

- read data from a file
- similar to rufs\_write.
- we made sure to only start reading from the block containing the offset and stopped at the block containing offset+size